

project_report

June 1, 2023

1 Monte Carlo Methods Project

1.1 Getting started

To run the code, first install Julia:

1. Visit <https://julialang.org/downloads/> and select the installer for your system (I used 1.8.3, but probably 1.9.0 will work just as well). If you're on Linux, run:

```
wget https://julialang-s3.julialang.org/bin/linux/x64/1.8/julia-1.8.3-linux-x86_64.tar.gz
tar zxvf julia-1.8.3-linux-x86_64.tar.gz
```

2. On Linux, that's all you have to do. You can choose to make the `julia` executable more accessible by, e.g., opening `~/.bashrc` and adding the line:

```
export PATH="$PATH:/path/to/<Julia directory>/bin"
```

3. Make sure to start a new terminal or run `source ~/.bashrc` afterwards.

Consequently, you want to install the relevant packages:

1. Move one level above the current directory (the current directory is the one that contains this file) and activate Julia:

```
cd ..
julia
```

2. Inside Julia, change to the Pkg mode by typing `]`. In there, you want to activate the project `toml` file and instantiate the project. This will install all necessary packages and precompile them.

```
]  
activate Project  
instantiate
```

```
[ ]: using LinearAlgebra  
      using Distributions  
      using Plots
```

General note: It was not quite clear to me whether the starting values should be returned by the sampling algorithms, or discarded. Either way, I don't think it matters much, but I decided to be prudent and discard the starting values.

2 Exercise (a)

```
[ ]: function quadraticpotential_grad(q)
    q
end

function leapfrog(q0, p0, tau, l, mass, potential_grad::Function)
    # Determine dimensionality
    d = length(q0)

    # Also accomodate initial values
    qs = [Vector{Float64}(undef, d) for _ = 1:(l+1)]
    ps = [Vector{Float64}(undef, d) for _ = 1:(l+1)]

    qs[1] = q0
    ps[1] = p0

    for i = 2:(l+1)
        p_half = ps[i - 1] .- tau / 2 .* potential_grad(qs[i - 1])
        qs[i] = qs[i - 1] .+ tau * inv(mass) * p_half
        ps[i] = p_half .- tau / 2 .* potential_grad(qs[i])
    end

    # Discard initial value
    (qs[2:(l+1)], ps[2:(l+1)])
end
```

leapfrog (generic function with 1 method)

3 Exercise (b)

```
[ ]: q0 = [0.]
      p0 = [1.]
      tau = 1.2
      l = 20
      mass = [1.;;]

      (ps1, qs1) = leapfrog(q0, p0, tau, l, mass, quadraticpotential_grad)

      tau = 0.3
      l = 80


      (ps2, qs2) = leapfrog(q0, p0, tau, l, mass, quadraticpotential_grad);

[ ]: qs1_flat = [v[1] for v in qs1]
      ps1_flat = [v[1] for v in ps1]
```


```
qs2_flat = [v[1] for v in qs2]
ps2_flat = [v[1] for v in ps2];
```

```
[ ]: t = collect(0:0.01:2)
q_exact = -cos.(t .+ /2)
p_exact = sin.(t .+ /2)

plot(q_exact, p_exact)
plot!(qs1_flat, ps1_flat)
```



```
[ ]: plot(q_exact, p_exact)
plot!(qs2_flat, ps2_flat)
```



4 Exercise (c)

Substituting Equation (4) into (7), and then using Equation (8), we get:

$$\begin{aligned}
 \pi(\mathbf{q}, \mathbf{p}) &\propto \exp(-H(\mathbf{q}, \mathbf{p})) \\
 &\propto e^{-U(\mathbf{q})-K(\mathbf{p})} \\
 &\propto e^{-\ln(\pi(\mathbf{q}))} e^{-\frac{1}{2}\mathbf{p}^T M^{-1} \mathbf{p}} \\
 &\propto \pi(\mathbf{q}) e^{-\frac{1}{2}\mathbf{p}^T M^{-1} \mathbf{p}}.
 \end{aligned}$$

If we fix \mathbf{p} , we can treat it as a constant:

$$\begin{aligned}
 \pi(\mathbf{q} \mid \mathbf{p}) &\propto \pi(\mathbf{q}) e^{-\frac{1}{2}\mathbf{p}^T M^{-1} \mathbf{p}} \\
 &\propto \pi(\mathbf{q}).
 \end{aligned}$$

This shows that sampling from $\pi(\mathbf{q}, \mathbf{p})$ is equivalent to sampling from $\pi(\mathbf{q})$ if \mathbf{p} is fixed.

5 Exercise (d)

Fixing \mathbf{q} to specify a pdf for \mathbf{p} yields:

$$\begin{aligned}
\pi(\mathbf{p} \mid \mathbf{q}) &\propto \pi(\mathbf{q}) e^{-\frac{1}{2} \mathbf{p}^T M^{-1} \mathbf{p}} \\
&\propto e^{-\frac{1}{2} \mathbf{p}^T M^{-1} \mathbf{p}} \\
&\propto (2\pi)^{-d/2} \det(M)^{-1/2} e^{-\frac{1}{2} \mathbf{p}^T M^{-1} \mathbf{p}}.
\end{aligned}$$

In other words, $\pi(\mathbf{p} \mid \mathbf{q})$ is equal to a 0-mean multivariate Gaussian pdf up to a normalising constant. M can be interpreted as the covariance matrix of that Gaussian.

6 Exercise (e)

$-H(\mathbf{q}^*, \mathbf{p}^*) + H(\mathbf{q}, \mathbf{p})$ can be interpreted as $\frac{\Delta H}{\Delta t}$. We have seen above that with an exact trajectory, the Hamiltonian is conserved: $\frac{dH}{dt} = 0$. Under this condition, the acceptance probability becomes:

$$\begin{aligned}
\alpha &= \min \{1, \exp(-H(\mathbf{q}^*, \mathbf{p}^*) + H(\mathbf{q}, \mathbf{p}))\} \\
&= \min \{1, \exp(0)\} \\
&= 1.
\end{aligned}$$

This means that using an approach that approximates the trajectory well would lead to acceptance probabilities close to 1.

7 Exercise (f)

```
[ ]: function quadraticpotential(q)
    0.5 * q' * q
end

function hamiltonian(q, p, mass, potential::Function)
    potential(q) + 0.5 * p' * inv(mass) * p
end

function hmc(n, q0, tau_dist, l, potential::Function, potential_grad::Function)
    # Determine dimensionality
    d = length(q0)

    # Also accomodate initial values
    qs = [Vector{Float64}(undef, d) for _ = 1:(n+1)]
    ps = [Vector{Float64}(undef, d) for _ = 1:(n+1)]

    # Julia can infer the size of I from the mathematical context
    mass = I
    norm = MvNormal(zeros(Float64, d), mass)

    qs[1] = q0
```

```

for t = 2:(n+1)
    ps[t - 1] = vec(rand(norm, 1))

    tau = rand(tau_dist, 1)[1]
    # Note that qs_star has a different dimensionality (l) than qs (n)
    (qs_star, ps_star) = leapfrog(qs[t - 1], ps[t - 1], tau, l, mass, ↵
    ↵potential_grad)
    q_star = qs_star[l]
    p_star = ps_star[l]

    h_prev = hamiltonian(qs[t - 1], ps[t - 1], mass, potential)
    h_star = hamiltonian(q_star, p_star, mass, potential)

    alpha = min(1, exp(h_prev - h_star))

    if rand() < alpha
        qs[t] = q_star
        ps[t] = p_star
    else
        qs[t] = qs[t - 1]
        ps[t] = ps[t - 1]
    end
end

(qs[2:(n+1)], ps[2:(n+1)])
end

```

hmc (generic function with 1 method)

Testing with the quadratic potential and its gradient. Since we are not applying a -log to the potential, we are essentially sampling from:

$$\begin{aligned}
 \pi(\mathbf{q} \mid \mathbf{p}) &\propto e^{-U(\mathbf{q})-K(\mathbf{p})} \\
 &\propto e^{-U(\mathbf{q})} \\
 &\propto e^{-\frac{1}{2}\mathbf{q}^T\mathbf{q}},
 \end{aligned}$$

which is proportional to a multivariate normal pdf with mean 0 and covariance matrix I . In other words, we should get a normal distribution if we set $d = 1$, which the histogram below shows to be the case.


```

[ ]: tau_dist = Uniform(0.5, 1.5)
samples, _ = hmc(100000, q0, tau_dist, l, quadraticpotential, ↵
    ↵quadraticpotential_grad);

```

```
[ ]: samples_flat = [s[1] for s in samples]
      histogram(samples_flat)

      x = -5:0.1:5
      plot!(x, pdf.(Normal(0, 1), x) .* 1e4, linewidth = 3)
```



8 Exercise (g)

```
[ ]: function rw_mh(n, start, prop_dist, target::Function)
      d = length(start)

      samples = [Vector{Float64}(undef, d) for _ = 1:(n+1)]
      samples[1] = start

      for i = 2:(n+1)
          sample_prop = vec(rand(prop_dist, 1)) + samples[i - 1]

          dens_prev = target(samples[i - 1])
          dens_prop = target(sample_prop)
          alpha = min(1, dens_prop / dens_prev)

          if rand() < alpha
```

```

        samples[i] = sample_prop
    else
        samples[i] = samples[i - 1]
    end
end

samples[2:(n+1)]
end

```


rw_mh (generic function with 1 method)

```

[ ]: target(x) = exp(-quadraticpotential(x))
samples = rw_mh(1000000, [0.], MvNormal(zeros(1), I), target)
histogram([s[1] for s in samples])

x = -5:0.1:5
plot!(x, target.(x) .* 2e4, linewidth = 3)

```



9 Exercise (h)

```
[ ]: function rand_sigma_mh(n, start, sigma_dist, target::Function)
    mh_samples = [Vector{Float64}(undef, d) for _ = 1:(n+1)]
    mh_samples[1] = start

    mu = zeros(size(start))

    for i = 2:(n+1)
        sigma = rand(sigma_dist, 1)[1]
        prop_dist = MvNormal(mu, sigma * I)

        mh_samples[i] = rw_mh(1, mh_samples[i - 1], prop_dist, target)[1]
    end

    mh_samples[2:(n+1)]
end
```

rand_sigma_mh (generic function with 1 method)

To be able to use Hamiltonian Monte Carlo, we first need to find the gradient $\nabla_{\mathbf{q}} U$. Since

$$\nabla_{\mathbf{q}} U = \begin{pmatrix} \frac{\partial U}{\partial q_1} \\ \vdots \\ \frac{\partial U}{\partial q_d} \end{pmatrix},$$

we need to find $\frac{\partial U}{\partial q_i}$, where $i = 1, \dots, d$. Because we want to sample from a multivariate normal, we get the -log of its pdf:

$$\begin{aligned} \frac{\partial U}{\partial q_i} &= \frac{\partial}{\partial q_i} \left[-\log \left((2\pi)^{-\frac{d}{2}} \det(\Sigma)^{-\frac{1}{2}} e^{-\frac{1}{2} \mathbf{q}^T \Sigma^{-1} \mathbf{q}} \right) \right] \\ &= \frac{\partial}{\partial q_i} \left[-\left(\log \left((2\pi)^{-\frac{d}{2}} \det(\Sigma)^{-\frac{1}{2}} \right) + \log \left(e^{-\frac{1}{2} \mathbf{q}^T \Sigma^{-1} \mathbf{q}} \right) \right) \right] \\ &= \frac{\partial}{\partial q_i} \left[\frac{1}{2} \mathbf{q}^T \Sigma^{-1} \mathbf{q} \right]. \end{aligned}$$

Σ is diagonal, so its inverse can be found by taking the reciprocals of its diagonal values:

$$\begin{aligned}
\frac{\partial U}{\partial q_i} &= \frac{\partial}{\partial q_i} \left[\frac{1}{2} \mathbf{q}^T \begin{pmatrix} \frac{1}{\sigma_1} & \cdots & 0 \\ \vdots & \ddots & \vdots \\ 0 & \cdots & \frac{1}{\sigma_d} \end{pmatrix} \mathbf{q} \right] \\
&= \frac{\partial}{\partial q_i} \left[\frac{1}{2} \mathbf{q}^T \begin{pmatrix} \frac{1}{\sigma_1} q_1 \\ \vdots \\ \frac{1}{\sigma_d} q_d \end{pmatrix} \right] \\
&= \frac{\partial}{\partial q_i} \left[\frac{1}{2} \left(\frac{1}{\sigma_1} q_1^2 + \cdots + q_d^2 \right) \right] \\
&= \frac{1}{\sigma_i} q_i.
\end{aligned}$$

```
[ ]: sigmas = collect(0.01:0.01:1)
Σ = Diagonal(sigmas)
d = size(Σ, 1)
    = zeros(d)
target(x) = pdf(MvNormal( , Σ), x)

n = 1000
l = 150
tau_dist = Uniform(0.0104, 0.0156)
potential(x) = -log(target(x))
potential_grad(x) = @. (1 / sigmas) * x
samples_hmc, _ = hmc(n, , tau_dist, l, potential, potential_grad)

n = 150000
sigma_dist = Uniform(0.0176, 0.0264)
samples_mh = rand_sigma_mh(n, , sigma_dist, target)
samples_mh_thinned = samples_mh[1:150:length(samples_mh)];
```

Plotting the HMC and thinned MH samples (first on top of each other, then separately) shows that the MH chain moves/mixes much slower than the HMC one. In other words, there seems to be more autocorrelation in the MH samples.


```
[ ]: last_comp_hmc = [v[d] for v in samples_hmc]
last_comp_mh_th = [v[d] for v in samples_mh_thinned]

plot(last_comp_hmc, labels = "HMC")
p1 = plot!(last_comp_mh_th, labels = "MH (thinned)")

p2 = plot(last_comp_hmc, legend = false)

p3 = plot(last_comp_mh_th, color = 2, legend = false)

plot(p1, p2, p3, layout = (3, 1))
```



Plotting the sample variation over the component standard deviations, we see a fairly constant error for the HMC sampler (except for a little dip at the very end), while for the MH sampler, the error seems to increase with the standard deviation.

```
[ ]: # Turns the vector of vectors inside-out (components outside, samples inside)
comps_hmc = [[v[i] for v in samples_hmc] for i = 1:d]
means_hmc = mean.(comps_hmc)
plot(sigmas, means_hmc, labels = "HMC")

comps_mh = [[v[i] for v in samples_mh] for i = 1:d]
means_mh = mean.(comps_mh)
p1 = plot!(sigmas, means_mh, labels = "MH (thinned)")

p2 = plot(sigmas, means_hmc, legend = false)

p3 = plot(sigmas, means_mh, color = 2, legend = false)

plot(p1, p2, p3, layout = (3, 1))
```

